

Copiamos la traza pcap y el fichero hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar al cluster

```
merve@onur-ideacenter:~/Desktop/Infraestructura_para_big_data/3Evaluación_de_prestaciones$ scp traza_serde2.pcap uambd20@150.244.65.34:.
uambd20@150.244.65.34's password:
traza_serde2.pcap 100% 146MB 10.1MB/s 00:14
merve@onur-ideacenter:~/Desktop/Infraestructura_para_big_data/3Evaluación_de_prestaciones$ scp hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar uambd20@150.244.65.34:.
uambd20@150.244.65.34's password:
hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar 100% 3468KB 8.6MB/s 00:00

merve@onur-ideacenter:~/Desktop/Infraestructura_para_big_data/3Evaluación_de_prestaciones$ ssh uambd20@150.244.65.34
uambd20@150.244.65.34's password:
Last login: Wed Jan 31 00:14:25 2024 from 172.30.130.89
[uambd20@master ~]$ ls
hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar  traza_serde2.pcap  veryBig.txt
TestDFSIO_results.log                                veryBig2.txt
```

Creamos directorios libs y pcaps en el HDFS y movemos la traza pcap y el fichero jar

```
uambd20@master:~
File Edit View Search Terminal Help
mkdir: `/home/uambd20/pcaps': No such file or directory
[uambd20@master ~]$ hadoop fs -mkdir /user/uambd20/pcaps
[uambd20@master ~]$ hadoop fs -mkdir /user/uambd20/libs
[uambd20@master ~]$ hadoop fs -put traza_serde2.pcap /user/uambd20/pcaps
[uambd20@master ~]$ hadoop fs -put hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar /user/uambd20/libs
[uambd20@master ~]$ hadoop fs -ls /user/uambd20/pcaps
Found 1 items
-rw-r--r-- 3 uambd20 uambd20 153403944 2024-01-31 00:44 /user/uambd20/pcaps/traza_serde2.pcap
[uambd20@master ~]$ hadoop fs -ls /user/uambd20/libs
Found 1 items
-rw-r--r-- 3 uambd20 uambd20 3551149 2024-01-31 00:45 /user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
[uambd20@master ~]$
```

También creamos el directorio hive\_script para guardar los scripts de Hive

```
[uambd20@master ~]$ mkdir /home/uambd20/hive_scripts/tables
[uambd20@master ~]$ mkdir /home/uambd20/hive_scripts/otros
[uambd20@master ~]$ mkdir /home/uambd20/hive_scripts/ejercicios
[uambd20@master ~]$ ls
hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar  traza_serde2.pcap
hive_scripts                                              veryBig2.txt
TestDFSIO_results.log                                    veryBig.txt
[uambd20@master ~]$ cd hive_scripts/
[uambd20@master hive_scripts]$ ls
ejercicios  otros  tables
```

Por otro lado creamos los directorios en hive\_scripts para los scripts de tablas y ejercicios

Creamos una tabla pcaps para SerDe y adjunto a esta práctica el ficheros “create\_table\_serde.hql” correspondiente al que he llamado.

```
[uambd20@master hive scripts]$ cd tables/
[uambd20@master tables]$ ls
[uambd20@master tables]$ nano create_table_serde.hql
[uambd20@master tables]$ hive -f create_table_serde.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/10e463e3-3a52-496f-a63d-9f19d9b5810a_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to class path
Added resources: [hdfs:///user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.575 seconds
OK
Time taken: 0.111 seconds
OK
Time taken: 0.077 seconds
OK
Time taken: 0.139 seconds
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uambd20@master tables]$
```

Creamos un script para visualizar dirección IP origen(src) y número de bytes con protocol=”UDP”, también adjunto a esta práctica el ficheros “select\_udp.hql” correspondiente al que he llamado.

```
[uambd20@master otros]$ cd ..
[uambd20@master hive scripts]$ cd tables/
[uambd20@master tables]$ cd ..
[uambd20@master hive scripts]$ cd otros/
[uambd20@master otros]$ nano select_udp.hql
[uambd20@master otros]$ hive -f select_udp.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/ff31a7f3-3aaa-47cd-9b7d-78a714e24db2_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to class path
Added resources: [hdfs:///user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.499 seconds
Query ID = uambd20_20240131135454_21d74027-1828-4ae4-9de2-3d54ecc4241b
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1706529435081_0084, Tracking URL = http://master:8088/proxy/application_1706529435081_0084/
Kill Command = /opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/lib/hadoop/bin/hadoop job -kill job_1706529435081_0084
```

```
Ended Job = job_1706529435081_0085
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 3 Cumulative CPU: 16.67 sec HDFS Read: 306827971 HDFS Write: 569 SUCCESS
Stage-Stage-2: Map: 3 Reduce: 1 Cumulative CPU: 5.6 sec HDFS Read: 9567 HDFS Write: 163 SUCCESS
Total MapReduce CPU Time Spent: 22 seconds 270 msec
OK
91.189.91.157 96
91.189.89.198 96
fe80:0:0:250:56ff:fec0:8 354
192.168.158.254 1200
0.0.0.0 1200
192.168.158.1 11860
192.168.158.128 260352
192.168.158.2 970018
Time taken: 42.76 seconds, Fetched: 8 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uambd20@master otros]$
```

Creamos una tabla http y adjunto a esta práctica el ficheros “create\_http\_table.hql” correspondiente al que he llamado.

```
GNU nano 2.3.1 File: create http table.hql
ADD JAR hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdb20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.HttpPcapReader;
CREATE DATABASE IF NOT EXISTS uamdb20;
USE uamdb20;
DROP TABLE http;
CREATE EXTERNAL TABLE http(ts bigint,protocol string,src string,src_port string,dst string, dst_port string,le$
ROW FORMAT SERDE 'net.ripe.hadoop.pcap.serde.PcapDeserializer'
STORED AS INPUTFORMAT 'net.ripe.hadoop.pcap.mrl.io.PcapInputFormat'
OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'
LOCATION 'hdfs://user/uamdb20/pcaps/';
```

```
uamdb20@master:~/hive_scripts/tables
File Edit View Search Terminal Help
[uamdb20@master tables]$ nano create http table.hql
[uamdb20@master tables]$ hive -f create_http_table.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive
-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/8e842628-e2d9-4a57-a0db-d2066d771455_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies
.jar] to class path
Added resources: [hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.564 seconds
OK
Time taken: 0.028 seconds
OK
Time taken: 0.008 seconds
OK
Time taken: 0.067 seconds
OK
Time taken: 0.143 seconds
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uamdb20@master tables]$
```

Creamos un script para visualizar líneas de respuesta HTTP con “UserAgent” cambiando el reader a la clase HTTP y adjunto a esta práctica el ficheros “http\_request.hql” correspondiente al que he llamado.

```
GNU nano 2.3.1 File: http request.hql
ADD JAR hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdb20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.HttpPcapReader;
SELECT http_request,`header_user-agent` FROM http where http_request is not NULL;
```

```
uamdb20@master:~/hive_scripts/otros
File Edit View Search Terminal Help
GET /css/rtve.2015/rtve.sections/rtve.rtve.css HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Gecko/2010
0101 Firefox/55.0
GET /pages/pfPortada/1.8.3/js/userFavor/UserFavor.js HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Ge
cko/20100101 Firefox/55.0
GET /vpaid/vPlayer/player/v8.2.235/assets/video-js/video.js HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:
55.0) Gecko/20100101 Firefox/55.0
GET /vpaid/units/17_2_1/assets/css/cm0sUnit.css HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Ge
cko/20100101 Firefox/55.0
GET /css?family=Roboto+Slab:300,400,500,100,100italic,300,300italic,400italic,500italic,700,700italic,900,900italic
 HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Gecko/20100101 Firefox/55.0
POST /gsalphasha2g2 HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Gecko/20100101 Firefox/55.0
GET /piwik/piwik.php?action_name=Nivel%201&idsite=1&rec=1&r=944933&h=3&m=35&s=58&url=http%3A%2F%2Fwww.madrid.org%2F
cs%2FSatellite%3Fpagename%3DComunidadMadrid%252FHome&urlref=http%3A%2F%2Fwww.madrid.org%2Fcs%2FSatellite%3Fcid%3D12
73078206797%26language%3Des%26pagename%3DComunidadMadrid%252FEstructura&_id=2ef494f588f51635&_idts=1511782494&_idvc
=1&_idn=0&_refsts=0&_views=1511782494&send_image=0&cookie=1&res=1916x958&gt_ms=26 HTTP/1.1 Mozilla/5.0 (X11; Ub
untu; Linux x86_64; rv:55.0) Gecko/20100101 Firefox/55.0
GET /vpaid/units/14_12_0/creatives/creative_js.js HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Ge
cko/20100101 Firefox/55.0
GET /vpaid/vPlayer/player/v8.2.235/0vaMediaPlayer.js HTTP/1.1 Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Ge
cko/20100101 Firefox/55.0
Time taken: 23.833 seconds, Fetched: 2414 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uamdb20@master otros]$
```

Creamos una tabla dns y adjunto a esta práctica el ficheros “create\_table\_dns.hql” correspondiente al que he llamado.

```
GNU nano 2.3.1 File: create table dns.hql Modified
ADD JAR hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.DnsPcapReader;
CREATE DATABASE IF NOT EXISTS uamdb20;
USE uamdb20;
DROP TABLE dns;
CREATE EXTERNAL TABLE dns(ts bigint, src string,dst string,dns_queryid int,dns_qr boolean,dns_opcode string,dns_rcode string, dns_qname string,$
ROW FORMAT SERDE 'net.ripe.hadoop.pcap.serde.PcapDeserializer'
STORED AS INPUTFORMAT 'net.ripe.hadoop.pcap.mr1.io.PcapInputFormat'
OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'
LOCATION 'hdfs://user/uamdb20/pcaps/';
```

## EJERCICIOS:

**Ejercicio1 (0,75 puntos):** Creamos un script HIVE que obtiene la serie temporal en bits/s con granularidad 1 segundo. El script selecciona los tiempos y los bits, que se han convertido a bits multiplicando los bytes por 8.

```
uambd20@master:~/hive_scripts/ejercicios
File Edit View Search Terminal Help
GNU nano 2.3.1 File: ejercicio1.hql
ADD JAR hdfs:///user/uambd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uambd20;
SELECT ts,sum(len)*8 as bits FROM pcaps group by ts order by ts asc;

^G Get Help      ^O WriteOut      ^R Read File     ^Y Prev Page     ^K Cut Text       ^C Cur Pos
^X Exit          ^J Justify       ^W Where Is      ^V Next Page     ^U UnCut Text     ^T To Spell
```

```
1511782531      577744
1511782532        0
1511782533      20624
1511782534       496
1511782535      181728
1511782536       1472
1511782537        0
1511782538     2538880
1511782539      544512
1511782540     2247504
1511782541      531744
1511782542     396912
1511782543     569504
1511782544     272336
1511782545     40352
1511782546     47440
1511782547      5632
1511782548     139392
1511782549     2135392
1511782550     189280
1511782551     1570176
1511782552     1849536
1511782553     4371680
1511782554     1374656
1511782555     3127440
1511782556     156480
1511782557     65312
1511782558     26352
1511782559     57632
1511782560     17376
1511782561      992
1511782562     2944
1511782563     22224
1511782564     36880
1511782565      992
1511782566        0
1511782567     9824
1511782568     9040
1511782569     12672
1511782570     12944
1511782571     4497168
1511782572     18320
1511782573     2944
Time taken: 42.174 seconds, Fetched: 894 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uambd20@master ejercicios]$
```



**Ejercicio2 (0,75 puntos):** Creamos un script HIVE que obtiene registros de flujos junto con el número de bytes y paquetes. Los registros de flujos incluyen información sobre tiempo mínimo, tiempo máximo, dirección IP origen (src) y dirección IP destino (dst), puerto origen (src\_port) y puerto destino (dst\_port), protocolo (TCP o UDP), número de bytes (sum(len)), número de paquetes.

```
GNU nano 2.3.1 File: ejercicio2.hql
ADD JAR hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdbd20;
SELECT min(ts) as first_packet, max(ts) as last_packet, src, src_port, dst, dst_port, protocol, sum(len) as total_bytes, count(*) as total_packets
FROM pcaps
WHERE src is not NULL and dst is NOT NULL
GROUP BY src, src_port, dst, dst_port, protocol
ORDER BY total_bytes DESC;
```

```
[uamdbd20@master ejercicios]$ nano ejercicio2.hql
[uamdbd20@master ejercicios]$ time hive -f ejercicio2.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/c216b23f-4fb1-4d8f-92c7-594e2423e336_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to class path
Added resources: [hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.578 seconds
Query ID = uamdbd20_20240131145555_7c6752df-97cb-403c-8cb3-0ca3c7637cb9
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1706529435081_0094, Tracking URL = http://master:8088/proxy/application_1706529435081_0094/
```

uambd20@master:~/hive\_scripts/ejercicios

File Edit View Search Terminal Help

1511782000	1511782005	192.168.158.128	46962	130.206.192.10	80	TCP	0	8
1511781874	1511781880	192.168.158.128	52670	216.58.210.129	80	TCP	0	8
1511781711	1511781711	93.184.220.70	443	192.168.158.128	45768	TCP	0	6
1511782220	1511782225	192.168.158.128	45858	192.0.76.3	80	TCP	0	8
1511782556	1511782562	192.168.158.128	45066	8.254.94.254	80	TCP	0	8
1511782361	1511782366	192.168.158.128	45812	62.73.184.139	80	TCP	0	8
1511781892	1511781897	192.168.158.128	53674	67.226.210.13	80	TCP	0	8
1511781869	1511781869	192.168.158.128	52424	69.172.216.111	80	TCP	0	4
1511782355	1511782360	192.168.158.128	45758	62.73.184.139	80	TCP	0	8
1511782556	1511782556	192.168.158.128	45722	69.172.216.111	443	TCP	0	4
1511782516	1511782525	93.184.221.133	80	192.168.158.128	46894	TCP	0	6
1511781896	1511781896	192.168.158.128	45622	62.67.193.31	80	TCP	0	4
1511781892	1511781897	192.168.158.128	52192	2.18.115.117	80	TCP	0	8
1511781893	1511781898	192.168.158.128	47070	52.51.73.34	80	TCP	0	8
1511781896	1511781896	192.168.158.128	45624	62.67.193.31	80	TCP	0	4
1511781879	1511781885	192.168.158.128	52734	34.204.227.165	80	TCP	0	8
1511781996	1511781996	104.121.156.80	80	192.168.158.128	47160	TCP	0	2
1511781699	1511781704	192.168.158.128	52712	150.244.214.206	80	TCP	0	8
1511782183	1511782188	192.168.158.128	45462	216.58.214.174	80	TCP	0	8
1511782427	1511782432	4.23.62.126	80	192.168.158.128	47348	TCP	0	6
1511781890	1511781896	192.168.158.128	51334	62.67.193.21	80	TCP	0	8
1511781701	1511781706	150.244.214.21	80	192.168.158.128	42738	TCP	0	6
1511782303	1511782308	192.168.158.128	51242	216.58.201.130	80	TCP	0	8
1511782171	1511782176	192.168.158.128	52586	54.231.114.187	80	TCP	0	8
1511781879	1511781884	46.228.164.11	80	192.168.158.128	53942	TCP	0	6
1511782549	1511782554	216.58.211.194	80	192.168.158.128	53536	TCP	0	6
1511781855	1511781855	192.168.158.128	45154	52.52.228.19	443	TCP	0	2
1511781854	1511781854	192.168.158.128	45124	52.52.228.19	443	TCP	0	2
1511782419	1511782482	192.168.158.128	51148	217.148.69.24	443	TCP	0	14

Time taken: 41.161 seconds, Fetched: 8681 row(s)

WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.

WARN: Please see <http://www.slf4j.org/codes.html#release> for an explanation.

real 0m50.587s

user 0m21.728s

sys 0m1.936s

[uambd20@master ejercicios]\$

**Ejercicio3 (0,75 puntos):** Creamos un script HIVE que obtiene los ('header\_host') 'Host' (a nivel HTTP) más populares (en número de peticiones).

```
GNU nano 2.3.1                               File: ejercicio3.hql

ADD JAR hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdbd20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.HttpPcapReader;
SELECT `header_host`, count(*) as npet FROM http where http_request is not NULL group by `header_host` sort by npet asc;
```

```
[uamdbd20@master hive_scripts]$ cd ejercicios/
[uamdbd20@master ejercicios]$ ls
ejercicio1.hql  ejercicio2.hql
[uamdbd20@master ejercicios]$ nano ejercicio3.hql
[uamdbd20@master ejercicios]$ nano ejercicio3.hql
[uamdbd20@master ejercicios]$ hive -f ejercicio3.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/86cd9117-6ace-474c-9419-40029da8b1a4_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to class path
Added resources: [hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.542 seconds
```



```

adserver-us.adtech.advertising.com      12
bilder.bild.de      12
www.ucm.es          12
dle.rae.es          12
jadserve.postrelease.com      12
img.irtve.es        12
adserver.adtechus.com  12
gn.symcd.com        12
hackersonlineclub.com  12
aax-eu.amazon-adsystem.com    14
r.nexac.com          14
cdn.cnn.com          14
www.uam.es           14
active.cache.el-mundo.net      14
tags.bluekai.com      14
dt.adsafeprotected.com  14
dx.steelhousemedia.com  14
www.latimes.com       14
as.casalemedia.com     14
www.eluniversohoy.net  16
clients1.google.com    16
pixel.adsafeprotected.com      16
www.hpcn-uam.es       16
www.trbimg.com         16
cdn02.nintendo-europe.com      16
edition.i.cdn.cnn.com  18
tags.tiqcdn.com        18
ministeriodeagricultura.solution.weborama.fr  18
e00-elmundo.uecdn.es    18
ocsp.usertrust.com      20
aa.agkn.com             20
e00-marca.uecdn.es      20
ocsp.pki.goog           20
aax.amazon-adsystem.com  24
www.irtve.es           24
img.rtve.es            24
tribune-d.openx.net     28
ams1-ib.adnxs.com       30
beacon.krxd.net         30
secure.adnxs.com        46
ocsp.godaddy.com        48
ib.adnxs.com            56
ocsp.comodoca.com       72
fastlane.rubiconproject.com    100
Time taken: 49.822 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uamdb20@master ejercicios]$

```

**Ejercicio4 (0,75 puntos):** Creamos un script HIVE que obtiene la cantidad de 'UserAgent' (`header\_user-agent`) distintos junto con el número de peticiones observadas.

```
GNU nano 2.3.1 File: ejercicio4.hql
ADD JAR hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdb20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.HttpPcapReader;
SELECT `header_user-agent`, count(*) as npet FROM http where http_request is not NULL group by `header_user-agent` sort by npet asc;
```

```
[uamdb20@master ejercicios]$ ls
ejercicio1.hql ejercicio2.hql ejercicio3.hql
[uamdb20@master ejercicios]$ nano ejercicio4.hql
[uamdb20@master ejercicios]$ hive -f ejercicio4.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.
1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/98492c04-d7af-4ce1-a0b3-fd96d7c35d5e_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to c
lass path
Added resources: [hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.581 seconds
Query ID = uamdb20_20240131162626_62570f37-ffbb-496b-abb9-5e2ede5447fe
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
```

```
uamdb20@master:~/hive_scripts/ejercicios
File Edit View Search Terminal Help
2024-01-31 16:27:18,797 Stage-2 map = 0%, reduce = 0%
2024-01-31 16:27:24,034 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.36 sec
2024-01-31 16:27:30,317 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.14 sec
MapReduce Total cumulative CPU time: 3 seconds 140 msec
Ended Job = job_1706529435081_0114
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 3 Cumulative CPU: 36.03 sec HDFS Read: 306828779 HDFS Write: 385 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 3.14 sec HDFS Read: 5743 HDFS Write: 82 SUCCESS
Total MapReduce CPU Time Spent: 39 seconds 170 msec
OK
Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Gecko/20100101 Firefox/55.0 2414
Time taken: 50.73 seconds, Fetched: 1 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uamdb20@master ejercicios]$
```

Solo se observa Mozilla/5.0 con un total de 2414 peticiones.

**Ejercicio 5 (0,75 puntos):** Creamos un script HIVE que obtiene las diferentes líneas de respuesta HTTP (http\_response) junto con el número de veces que se han observado.

```
GNU nano 2.3.1 File: ejercicio5.hql
Add JAR hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdbd20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.HttpPcapReader;
SELECT http_response, count(*) as nresp FROM http where http_response is not NULL group by http_response sort by nresp asc;
```

```
[uamdbd20@master ejercicios]$ nano ejercicio5.hql
[uamdbd20@master ejercicios]$ hive -f ejercicio5.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-
common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/d92cca07-17e7-4504-9e5f-e668b7eeb051_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.
jar] to class path
Added resources: [hdfs:///user/uamdbd20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.56 seconds
Query ID = uamdbd20_20240131163333_38001660-0722-4f74-bc7d-3cd98a1b9818
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1706529435081_0115, Tracking URL = http://master:8088/proxy/application_1706529435081_0115/
```

```
uamdbd20@master:~/hive_scripts/ejercicios

File Edit View Search Terminal Help
Ended Job = job_1706529435081_0116
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 3 Cumulative CPU: 36.09 sec HDFS Read: 306826605 HDFS Write: 1133 SUCCESS
Stage-Stage-2: Map: 3 Reduce: 1 Cumulative CPU: 5.52 sec HDFS Read: 10143 HDFS Write: 532 SUCCESS
Total MapReduce CPU Time Spent: 41 seconds 610 msec
OK
HTTP/1.1 400 Bad Request 2
HTTP/1.1 307 Temporary Redirect 2
HTTP/1.1 204 No Response 2
HTTP/1.1 302 nginx/1.11.3 2
HTTP/1.1 206 Partial Content 2
HTTP/1.1 204 2
HTTP/1.1 101 Switching Protocols 2
HTTP/1.0 301 Moved Permanently 2
HTTP/1.0 200 OK 2
HTTP/1.1 303 See Other 4
HTTP/1.1 404 Not Found 4
HTTP/1.1 302 Moved Permanently 8
HTTP/1.0 400 Bad request 8
HTTP/1.1 302 12
HTTP/1.1 200 14
HTTP/1.1 301 Moved Permanently 32
HTTP/1.1 204 No Content 60
HTTP/1.1 302 Moved Temporarily 96
HTTP/1.1 302 Found 250
HTTP/1.1 200 OK 1792
Time taken: 50.452 seconds, Fetched: 20 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uamdbd20@master ejercicios]$
```

**Ejercicio 6 (0,75 puntos):** Creamos un script HIVE que obtiene la popularidad de los nombre de dominio (dns\_qname) por número de peticiones.

```
GNU nano 2.3.1 File: ejercicio6.hql

ADD JAR hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar;
USE uamdb20;
SET net.ripe.hadoop.pcap.io.reader.class=net.ripe.hadoop.pcap.DnsPcapReader;
SELECT dns_qname, count(*) as requests
FROM dns
WHERE dns_qr=false
GROUP BY dns_qname
ORDER BY requests DESC;
```

```
[uamdb20@master hive scripts]$ cd ejercicios/
[uamdb20@master ejercicios]$ ls
ejercicio1.hql ejercicio2.hql ejercicio3.hql ejercicio4.hql ejercicio5.hql
[uamdb20@master ejercicios]$ nano ejercicio6.hql
[uamdb20@master ejercicios]$ hive -f ejercicio6.hql

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties
converting to local hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar
Added [/tmp/5923b8d9-429a-4c55-91a7-1832e5e0d1ac_resources/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar] to class path
Added resources: [hdfs:///user/uamdb20/libs/hadoop-pcap-serde-1.2-SNAPSHOT-jar-with-dependencies.jar]
OK
Time taken: 1.609 seconds
Query ID = uamdb20_20240131182020_634c9589-faba-4e4b-8f3d-006612482b0e
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1706529435081_0117, Tracking URL = http://master:8088/proxy/application_1706529435081_0117/
```

```
Total MapReduce CPU Time Spent: 46 seconds 560 msec
OK
bidder.criteo.com.      56
pix04.revsci.net.      48
www.facebook.com.      48
www.youtube.com.       46
www.google.com. 40
ib.adnxs.com. 38
securepubads.g.doubleclick.net. 38
connect.facebook.net. 36
versioncheck-bg.addons.mozilla.org. 36
api.permutive.com. 36
www.google-analytics.com. 36
adservice.google.com. 34
pagead2.googlesyndication.com. 34
syndication.twitter.com. 34
d2tjil5tz17cbz.cloudfront.net. 32
d2tjil5tz17cbz.cloudfront.net.localdomain. 32
ad.doubleclick.net. 32
www.googletag services.com. 32
rb.adnxs.com. 32
b.scorecardresearch.com. 32
platform.twitter.com. 30
as.casalemedia.com. 28
fonts.googleapis.com. 28
tpc.googlesyndication.com. 28
www.google.es. 28
ajax.googleapis.com. 26
staticxx.facebook.com. 24
ping.chartbeat.net. 24
pix04.revsci.net.localdomain. 24
cm.g.doubleclick.net. 24
secure.adnxs.com. 24
googleads.g.doubleclick.net. 24
beacon.krx.net. 24
bs.serving-sys.com. 24
acdn.adnxs.com. 24
pbs.twimg.com. 24
fonts.gstatic.com. 22
```

uambd20@master:~/hive\_scripts/ejercicios

File Edit View Search Terminal Help

```
asset.flite.com. 4
apps.uam.es. 4
api.pinterest.com. 4
api.elconfidencial.com. 4
api.bounceexchange.com. 4
analytics.brandcrumb.com. 4
an.pxl.ace.advertising.com. 4
amplify.outbrain.com. 4
amazon.es. 4
alumni.uam.es. 4
akl.abmr.net. 4
ak.sail-horizon.com. 4
ak-ads-ns.prisas.com. 4
advertising.amazon.com. 4
adserver.adtechus.com. 4
adserver-us.adtech.advertising.com. 4
ads.undertone.com. 4
ad.yieldlab.net. 4
a.visualrevenue.com. 4
a.rfihub.com. 4
a.amz.mshcdn.com. 4
a-ssl.ligatus.com. 4
7li.nuggad.net. 4
5808301.fl.doubleclick.net. 4
58011369.log.optimizely.com. 4
5365628.fl.doubleclick.net. 4
4711.lineadirecta.com. 4
4.bp.blogspot.com. 4
33across-d.openx.net. 4
1tol.bbva.com. 4
as.es. 2
quit. 2
local. 2
quit.localdomain. 2
ww.google.es. 2
as.com. 2
google.es. 2
Time taken: 52.571 seconds, Fetched: 1024 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[uambd20@master ejercicios]$
```

**Ejercicio 7 (1 punto):** Obtenemos los valores por defecto de los siguientes parámetros (desde HIVE) usando los comandos:

```
set mapreduce.map.memory.mb;
```

(La memoria (en MB) dedicada a las tareas del mapper es 0)

```
set mapreduce.map.java.opts;
```

(Las opciones de la JVM que se usarán para las tareas del mapper es TRUE)

```
set mapreduce.reduce.memory.mb;
```

(La memoria (en MB) dedicada a las tareas del reducer es 0)

```
set mapreduce.reduce.java.opts;
```

(Las opciones de la JVM que se usarán para las tareas del mapper es TRUE)

```
set mapreduce.task.io.sort.mb;
```

(La cantidad de memoria (en MB) usada para realizar ordenación de ficheros es 256)

```
set yarn.scheduler.minimum-allocation-mb;
```

(El valor mínimo de memoria (en MB) usado para un container es 1024)

```
set yarn.scheduler.maximum-allocation-mb;
```

(El valor máximo de memoria (en MB) usado para un container es 65536)

```
set yarn.scheduler.minimum-allocation-vcores;
```

(El número mínimo de vCores que puede usar un container es 1)

```
set yarn.scheduler.maximum-allocation-vcores;
```

(El número máximo de vCores que puede usar un container es 32)



```
uamdb20@master:~  
File Edit View Search Terminal Help  
merve@onur-ideacenter:~/Desktop/Infraestructura_para_big_data/3Evaluación_de_pre  
staciones$ ssh uamdb20@150.244.65.34  
uamdb20@150.244.65.34's password:  
Last login: Wed Jan 31 18:02:28 2024 from 172.30.130.89  
[uamdb20@master ~]$ hive  
  
Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.  
1-1.cdh5.15.1.p0.4/jars/hive-common-1.1.0-cdh5.15.1.jar!/hive-log4j.properties  
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.  
hive> set mapreduce.map.memory.mb;  
mapreduce.map.memory.mb=0  
hive> set mapreduce.map.java.opts;  
mapreduce.map.java.opts=-Djava.net.preferIPv4Stack=true  
hive> set mapreduce.reduce.memory.mb;  
mapreduce.reduce.memory.mb=0  
hive> set mapreduce.reduce.java.opts;  
mapreduce.reduce.java.opts=-Djava.net.preferIPv4Stack=true  
hive> set mapreduce.task.io.sort.mb;  
mapreduce.task.io.sort.mb=256  
hive> set yarn.scheduler.minimum-allocation-mb;  
yarn.scheduler.minimum-allocation-mb=1024  
hive> set yarn.scheduler.maximum-allocation-mb;  
yarn.scheduler.maximum-allocation-mb=65536  
hive> set yarn.scheduler.minimum-allocation-vcores;  
yarn.scheduler.minimum-allocation-vcores=1  
hive> set yarn.scheduler.maximum-allocation-vcores;  
yarn.scheduler.maximum-allocation-vcores=32  
hive> 
```

**Ejercicio 8 (1,5 puntos):** Utilizando el script desarrollado en ejercicio 2, observar la influencia del parámetro “mapreduce.map.cpu.vcores”. Variar su valor entre 1 y 3. Realizar 4 ejecuciones con cada valor y obtener la media y al desviación estándar. Representar los datos de forma gráfica o mediante una tabla.

```
hive> set mapreduce.map.cpu.vcores;  
mapreduce.map.cpu.vcores=1
```

El valor por defecto del parámetro mapreduce.map.cpu.vcores es 1.

Realizamos 4 ejecuciones del script de Hive cuando el parámetro mapreduce.map.cpu.vcores está configurado en 1.

```
[uamdb20@master ejercicios]$ time hive -f ejercicio2.hql --hiveconf mapreduce.map.cpu.vcores=1
```

```
1511782355      1511782360      192.168.158.128 55190      172.227.120.189 80      TCP      0      8  
1511781896      1511781896      192.168.158.128 44540      185.33.223.83  80      TCP      0      4  
Time taken: 42.247 seconds, Fetched: 8681 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real    0m51.606s  
user    0m21.511s  
sys     0m1.947s  
[uamdb20@master ejercicios]$
```

```
1511781854      1511781854      192.168.158.128 45124      52.52.228.19   443     TCP      0      2  
1511782419      1511782482      192.168.158.128 51148      217.148.69.24  443     TCP      0      14  
Time taken: 42.0 seconds, Fetched: 8681 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real    0m50.935s  
user    0m21.005s  
sys     0m1.931s  
[uamdb20@master ejercicios]$
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8  
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8  
Time taken: 40.883 seconds, Fetched: 8681 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real    0m49.743s  
user    0m20.324s  
sys     0m1.762s  
[uamdb20@master ejercicios]$
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8  
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8  
Time taken: 44.439 seconds, Fetched: 8681 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real    0m53.686s  
user    0m21.403s  
sys     0m2.110s
```

Realizamos 4 ejecuciones del script de Hive cuando el parámetro `mapreduce.map.cpu.vcores` está configurado en 2.

```
[uamdb20@master ejercicios]$ time hive -f ejercicio2.hql --hiveconf mapreduce.map.cpu.vcores=2
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 42.022 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m51.591s
user    0m22.069s
sys     0m1.959s
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 41.35 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.770s
user    0m21.998s
sys     0m1.846s
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 41.042 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.199s
user    0m21.202s
sys     0m1.929s
```

```
1511782355      1511782360      192.168.158.128 55190      172.227.120.189 80      TCP      0      8
1511781896      1511781896      192.168.158.128 44540      185.33.223.83 80      TCP      0      4
Time taken: 41.369 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.659s
user    0m22.060s
sys     0m1.952s
```

Realizamos 4 ejecuciones del script de Hive cuando el parámetro `mapreduce.map.cpu.vcores` está configurado en 3.

```
[uamdb20@master ejercicios]$ time hive -f ejercicio2.hql --hiveconf mapreduce.map.cpu.vcores=3
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 41.613 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.795s
user    0m21.094s
sys     0m2.010s
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 41.212 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.515s
user    0m22.039s
sys     0m1.872s
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 41.346 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.291s
user    0m20.934s
sys     0m1.871s
```

```
1511781902      1511781907      192.168.158.128 37794      54.245.233.125 80      TCP      0      8
1511782127      1511782132      192.168.158.128 42674      216.58.210.131 80      TCP      0      8
Time taken: 40.935 seconds, Fetched: 8681 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m50.221s
user    0m21.329s
sys     0m1.853s
```

Question 8	Mean of Execution Time	Standard Deviation of Execution Time
Real Time 1 mapreduce.map.cpu.vcore=1	51.492	1.431
Real Time 2 mapreduce.map.cpu.vcore=2	50.805	0.502
Real Time 3 mapreduce.map.cpu.vcore=3	50.456	0.224




Figure 1: Resultados de 4 ejecuciones con valores diferentes para un parámetro específico.

Realtime1,2 y 3 indica que mapreduce.map.cpu.vcore=1,2 y 3

Observamos en la gráfica que el valor óptimo para el parámetro mapreduce.map.cpu.vcores es 3. Porque se reduce el tiempo de ejecución medio y la desviación estándar del tiempo de ejecución.

#### En BoxPlot:

Línea dentro de la caja: Esta línea representa la mediana (Q2) del conjunto de datos. Da una idea de la tendencia central de los datos.

Bigotes: Los bigotes se extienden desde los bordes de la caja hasta los valores mínimo y máximo dentro de un rango determinado. El rango puede basarse en una distancia especificada desde el IQR o puede extenderse a los valores mínimo y máximo dentro de los datos que no se consideran valores atípicos.



**Ejercicio 9 (1,5 puntos):** Utilizando el script desarrollado en ejercicio 3, observar la influencia del parámetro “mapreduce.map.memory.mb”. Probar con los siguientes valores [2048,4096,8192] . Realizar 4 ejecuciones con cada valor y obtener la media y la desviación estándar. Representar los datos de forma gráfica o mediante una tabla. Al realizar las ejecuciones se deben tener en cuenta la influencia de otros parámetros del clúster como por ejemplo:

“mapreduce.map.java.opts”

“mapreduce.task.io.sort.mb”

Tomando en consideración las siguientes condiciones:

IO Sort Mem << Mapper Maximum JVM Heap

Mapper Maximum JVM Heap - IO Sort Mem > 768MB

Realizamos 4 ejecuciones con la configuración cuando mapreduce.map.memory.mb=2048 y los valores de mapreduce.map.java.opts, mapreduce.task.io.sort.mb(=256) son los valores por defecto.

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=2048
```

```
ib.adnxs.com      56
ocsp.comodoca.com  72
fastlane.rubiconproject.com  100
Time taken: 49.445 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m57.631s
user    0m20.862s
sys     0m1.747s
```

```
ib.adnxs.com      56
ocsp.comodoca.com  72
fastlane.rubiconproject.com  100
Time taken: 48.446 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m56.782s
user    0m20.490s
sys     0m1.586s
```

```
ib.adnxs.com      56
ocsp.comodoca.com  72
fastlane.rubiconproject.com  100
Time taken: 51.447 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m59.814s
user    0m20.284s
sys     0m1.738s
```

```
ib.adnxs.com      56
ocsp.comodoca.com  72
fastlane.rubiconproject.com  100
Time taken: 53.711 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    1m1.975s
user    0m21.585s
sys     0m1.603s
```



Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb=2048` y `mapreduce.map.java.opts` se establece en `-Xmx1000m` (Establece el tamaño máximo del montón en 1 GB) y `mapreduce.task.io.sort.mb` se establece en 100 (Asigna 100 MB de memoria para realizar ordenación de ficheros).

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=2048 --hiveconf mapreduce.map.java.opts=-Xmx1000m --hiveconf mapreduce.task.io.sort.mb=100
```

```
beacon.krxd.net 30
ams1-ib.adnxs.com 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 49.422 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m57.815s
user    0m21.169s
sys     0m1.499s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com 30
beacon.krxd.net 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 53.351 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    1m1.798s
user    0m21.143s
sys     0m1.710s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com 30
beacon.krxd.net 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 49.859 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m58.385s
user    0m21.434s
sys     0m1.700s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com 30
beacon.krxd.net 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 51.733 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m59.820s
user    0m20.634s
sys     0m1.718s
[uambd20@master ejercicios]$
```

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb=2048` y `mapreduce.map.java.opts` se establece en `-Xmx2000m` (Establece el tamaño máximo del montón en 2 GB) y `mapreduce.task.io.sort.mb` se establece en 300 (Asigna 300 MB de memoria para realizar ordenación de ficheros).

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=2048
--hiveconf mapreduce.map.java.opts=-Xmx2000m --hiveconf mapreduce.task.io.sort.mb=300
```

```
beacon.krxd.net 30
ams1-ib.adnxs.com 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 50.092 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m57.996s
user    0m20.771s
sys     0m1.530s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com 30
beacon.krxd.net 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 53.119 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    1m1.413s
user    0m20.846s
sys     0m1.642s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com 30
beacon.krxd.net 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 50.561 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m59.026s
user    0m21.881s
sys     0m1.743s
[uambd20@master ejercicios]$
```

```
beacon.krxd.net 30
ams1-ib.adnxs.com 30
secure.adnxs.com 46
ocsp.godaddy.com 48
ib.adnxs.com 56
ocsp.comodoca.com 72
fastlane.rubiconproject.com 100
Time taken: 47.963 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m56.469s
user    0m21.663s
sys     0m1.576s
[uambd20@master ejercicios]$
```

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb=4096` y los valores de `mapreduce.map.java.opts`, `mapreduce.task.io.sort.mb(=256)` son los valores por defecto.

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=4096
```

```
tribune-d.openx.net      28
ams1-ib.adnxs.com        30
beacon.krxd.net 30
secure.adnxs.com         46
ocsp.godaddy.com         48
ib.adnxs.com             56
ocsp.comodoca.com        72
fastlane.rubiconproject.com 100
Time taken: 50.718 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release(
) was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m58.851s
user    0m21.012s
sys     0m1.671s
[uambd20@master ejercicios]$
```

```
tribune-d.openx.net      28
ams1-ib.adnxs.com        30
beacon.krxd.net 30
secure.adnxs.com         46
ocsp.godaddy.com         48
ib.adnxs.com             56
ocsp.comodoca.com        72
fastlane.rubiconproject.com 100
Time taken: 51.894 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release(
) was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    1m0.426s
user    0m21.996s
sys     0m1.724s
[uambd20@master ejercicios]$
```

```
tribune-d.openx.net      28
ams1-ib.adnxs.com        30
beacon.krxd.net 30
secure.adnxs.com         46
ocsp.godaddy.com         48
ib.adnxs.com             56
ocsp.comodoca.com        72
fastlane.rubiconproject.com 100
Time taken: 49.435 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release(
) was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m57.813s
user    0m20.568s
sys     0m1.877s
[uambd20@master ejercicios]$
```

```
tribune-d.openx.net      28
beacon.krxd.net 30
ams1-ib.adnxs.com        30
secure.adnxs.com         46
ocsp.godaddy.com         48
ib.adnxs.com             56
ocsp.comodoca.com        72
fastlane.rubiconproject.com 100
Time taken: 48.257 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release(
) was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.

real    0m56.757s
user    0m20.950s
sys     0m1.747s
[uambd20@master ejercicios]$
```

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb= 4096` y `mapreduce.map.java.opts` se establece en `-Xmx1000m` (Establece el tamaño máximo del montón en 1 GB) y `mapreduce.task.io.sort.mb` se establece en 100 (Asigna 100 MB de memoria para realizar ordenación de ficheros).

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=4096 --hiveconf mapreduce.map.java.opts=-Xmx1000m --hiveconf mapreduce.task.io.sort.mb=100
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 46.043 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#
release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explana
tion.

real    0m54.332s
user    0m20.865s
sys     0m1.764s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 53.075 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#
release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explana
tion.

real    1m1.658s
user    0m21.537s
sys     0m1.597s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 54.515 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#
release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explana
tion.

real    1m2.798s
user    0m20.216s
sys     0m1.791s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 51.249 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#
release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explana
tion.

real    0m59.778s
user    0m21.974s
sys     0m1.819s
[uambd20@master ejercicios]$
```

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb= 4096` y `mapreduce.map.java.opts` se establece en `-Xmx2000m` (Establece el tamaño máximo del montón en 2 GB) y `mapreduce.task.io.sort.mb` se establece en 300 (Asigna 300 MB de memoria para realizar ordenación de ficheros).

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=4096
--hiveconf mapreduce.map.java.opts=-Xmx2000m --hiveconf mapreduce.task.io.sort.mb=300
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 52.027 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactor
y#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an expla
nation.

real    1m0.415s
user    0m21.080s
sys     0m1.770s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 57.793 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactor
y#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an expla
nation.

real    1m6.530s
user    0m20.926s
sys     0m1.691s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 50.589 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactor
y#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an expla
nation.

real    0m58.813s
user    0m20.696s
sys     0m1.603s
[uambd20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com      72
fastlane.rubiconproject.com 100
Time taken: 55.639 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactor
y#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an expla
nation.

real    1m4.208s
user    0m22.436s
sys     0m1.702s
[uambd20@master ejercicios]$
```

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb=8192` y los valores de `mapreduce.map.java.opts`, `mapreduce.task.io.sort.mb(=256)` son los valores por defecto.

```
[uamdb20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=8192
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com       72
fastlane.rubiconproject.com 100
Time taken: 46.137 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation
.
real    0m54.497s
user    0m20.688s
sys     0m1.726s
[uamdb20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com       72
fastlane.rubiconproject.com 100
Time taken: 50.501 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation
.
real    0m58.731s
user    0m20.822s
sys     0m1.822s
[uamdb20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com       72
fastlane.rubiconproject.com 100
Time taken: 48.197 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation
.
real    0m56.650s
user    0m21.762s
sys     0m1.599s
[uamdb20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com       46
ocsp.godaddy.com       48
ib.adnxs.com           56
ocsp.comodoca.com       72
fastlane.rubiconproject.com 100
Time taken: 49.555 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation
.
real    0m58.009s
user    0m21.471s
sys     0m1.676s
[uamdb20@master ejercicios]$
```



Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb= 8192` y `mapreduce.map.java.opts` se establece en `-Xmx1000m` (Establece el tamaño máximo del montón en 1 GB) y `mapreduce.task.io.sort.mb` se establece en 100 (Asigna 100 MB de memoria para realizar ordenación de ficheros).

```
[uambd20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=8192
--hiveconf mapreduce.map.java.opts=-Xmx1000m --hiveconf mapreduce.task.io.sort.mb=100
```

<pre>ams1-ib.adnxs.com      30 beacon.krxd.net 30 secure.adnxs.com      46 ocsp.godaddy.com      48 ib.adnxs.com 56 ocsp.comodoca.com      72 fastlane.rubiconproject.com 100 Time taken: 45.558 seconds, Fetched: 468 row(s) WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory# release() was invoked. WARN: Please see http://www.slf4j.org/codes.html#release for an explana tion.  real    0m54.061s user    0m20.042s sys     0m1.623s [uambd20@master ejercicios]\$</pre>	<pre>ams1-ib.adnxs.com      30 beacon.krxd.net 30 secure.adnxs.com      46 ocsp.godaddy.com      48 ib.adnxs.com 56 ocsp.comodoca.com      72 fastlane.rubiconproject.com 100 Time taken: 46.7 seconds, Fetched: 468 row(s) WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory# release() was invoked. WARN: Please see http://www.slf4j.org/codes.html#release for an explana tion.  real    0m55.396s user    0m21.062s sys     0m1.706s [uambd20@master ejercicios]\$</pre>
------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

<pre>beacon.krxd.net 30 ams1-ib.adnxs.com      30 secure.adnxs.com      46 ocsp.godaddy.com      48 ib.adnxs.com 56 ocsp.comodoca.com      72 fastlane.rubiconproject.com 100 Time taken: 48.433 seconds, Fetched: 468 row(s) WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory# release() was invoked. WARN: Please see http://www.slf4j.org/codes.html#release for an explana tion.  real    0m56.854s user    0m20.797s sys     0m1.669s [uambd20@master ejercicios]\$</pre>	<pre>ams1-ib.adnxs.com      30 beacon.krxd.net 30 secure.adnxs.com      46 ocsp.godaddy.com      48 ib.adnxs.com 56 ocsp.comodoca.com      72 fastlane.rubiconproject.com 100 Time taken: 45.132 seconds, Fetched: 468 row(s) WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory# release() was invoked. WARN: Please see http://www.slf4j.org/codes.html#release for an explana tion.  real    0m53.482s user    0m19.841s sys     0m1.755s [uambd20@master ejercicios]\$</pre>
------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Realizamos 4 ejecuciones con la configuración cuando `mapreduce.map.memory.mb= 8192` y `mapreduce.map.java.opts` se establece en `-Xmx2000m` (Establece el tamaño máximo del montón en 2 GB) y `mapreduce.task.io.sort.mb` se establece en 300 (Asigna 300 MB de memoria para realizar ordenación de ficheros).

```
[uamdb20@master ejercicios]$ time hive -f ejercicio3.hql --hiveconf mapreduce.map.memory.mb=8192
--hiveconf mapreduce.map.java.opts=-Xmx2000m --hiveconf mapreduce.task.io.sort.mb=300
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com      46
ocsp.godaddy.com      48
ib.adnxs.com          56
ocsp.comodoca.com     72
fastlane.rubiconproject.com 100
Time taken: 46.354 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory
#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explan
ation.

real    0m54.779s
user    0m20.321s
sys     0m1.749s
[uamdb20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com      46
ocsp.godaddy.com      48
ib.adnxs.com          56
ocsp.comodoca.com     72
fastlane.rubiconproject.com 100
Time taken: 49.466 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory
#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explan
ation.

real    0m58.308s
user    0m21.296s
sys     0m1.648s
[uamdb20@master ejercicios]$
```

```
ams1-ib.adnxs.com      30
beacon.krxd.net 30
secure.adnxs.com      46
ocsp.godaddy.com      48
ib.adnxs.com          56
ocsp.comodoca.com     72
fastlane.rubiconproject.com 100
Time taken: 46.349 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory
#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explan
ation.

real    0m54.736s
user    0m20.902s
sys     0m1.823s
[uamdb20@master ejercicios]$
```

```
beacon.krxd.net 30
ams1-ib.adnxs.com      30
secure.adnxs.com      46
ocsp.godaddy.com      48
ib.adnxs.com          56
ocsp.comodoca.com     72
fastlane.rubiconproject.com 100
Time taken: 49.393 seconds, Fetched: 468 row(s)
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory
#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explan
ation.

real    0m57.585s
user    0m20.883s
sys     0m1.602s
[uamdb20@master ejercicios]$
```

Question 9	Mean	Standard Deviation
Real Time 1 mapreduce.map.memory.mb=2048 mapreduce.map.java.opts=por defecto mapreduce.task.io.sort.mb=por defecto(256)	59.050	2.018
Real Time 2 mapreduce.map.memory.mb=2048 mapreduce.map.java.opts=-Xmx1000m mapreduce.task.io.sort.mb=100	59.454	1.538
Real Time 3 mapreduce.map.memory.mb=2048 mapreduce.map.java.opts=-Xmx2000m mapreduce.task.io.sort.mb=300	58.726	1.798
Real Time 4 mapreduce.map.memory.mb=4096 mapreduce.map.java.opts=por defecto mapreduce.task.io.sort.mb=por defecto(256)	58.462	1.354
Real Time 5 mapreduce.map.memory.mb=4096 mapreduce.map.java.opts=-Xmx1000m mapreduce.task.io.sort.mb=100	59.642	3.250
Real Time 6 mapreduce.map.memory.mb=4096 mapreduce.map.java.opts=-Xmx2000m mapreduce.task.io.sort.mb=300	62.491	3.045
Real Time 7 mapreduce.map.memory.mb=8192 mapreduce.map.java.opts=por defecto mapreduce.task.io.sort.mb=por defecto(256)	56.972	1.612
Real Time 8 mapreduce.map.memory.mb=8192 mapreduce.map.java.opts=-Xmx1000m mapreduce.task.io.sort.mb=100	54.948	1.301
Real Time 9 mapreduce.map.memory.mb=8192 mapreduce.map.java.opts=-Xmx2000m mapreduce.task.io.sort.mb=300	56.352	1.615




Figure 2: Resultados de 4 ejecuciones con valores diferentes para 3 parámetros distintos.

RealTime1 indica que `mapreduce.map.memory.mb=2048` `mapreduce.map.java.opts=por defecto`, `mapreduce.task.io.sort.mb=por defecto(256)`

RealTime2 indica que `mapreduce.map.memory.mb=2048`, `mapreduce.map.java.opts=-Xmx1000m`, `mapreduce.task.io.sort.mb=100`

RealTime3 indica que `mapreduce.map.memory.mb=2048`, `mapreduce.map.java.opts=-Xmx2000m`, `mapreduce.task.io.sort.mb=300`

RealTime4 indica que `mapreduce.map.memory.mb=4096` `mapreduce.map.java.opts=por defecto`, `mapreduce.task.io.sort.mb=por defecto(256)`

RealTime5 indica que `mapreduce.map.memory.mb=4096`, `mapreduce.map.java.opts=-Xmx1000m`, `mapreduce.task.io.sort.mb=100`

RealTime6 indica que `mapreduce.map.memory.mb=4096`, `mapreduce.map.java.opts=-Xmx2000m`, `mapreduce.task.io.sort.mb=300`

RealTime7 indica que `mapreduce.map.memory.mb=8192` `mapreduce.map.java.opts=por defecto`, `mapreduce.task.io.sort.mb=por defecto(256)`

RealTime8 indica que `mapreduce.map.memory.mb=8192`, `mapreduce.map.java.opts=-Xmx1000m`, `mapreduce.task.io.sort.mb=100`

RealTime9 indica que `mapreduce.map.memory.mb=8192`, `mapreduce.map.java.opts=-Xmx2000m`, `mapreduce.task.io.sort.mb=300`

Observamos en la gráfica que para reducir tanto el tiempo de ejecución medio como la desviación estándar del tiempo de ejecución, los valores óptimos de los parámetros son:

`mapreduce.map.memory.mb=8192`  
`mapreduce.map.java.opts=-Xmx1000m`  
`mapreduce.task.io.sort.mb=100.`

**Ejercicio 10 (1,5 puntos):** Utilizando el script desarrollado en ejercicio 6, observar la influencia de la aplicación de la vectorización de HIVE (activar todos los parámetros de vectorización) . Realizar 4 ejecuciones con cada una de las dos opciones y obtener la media y al desviación estándar. Representar los datos de forma gráfica o mediante una tabla.

Los parámetros de vectorización de HIVE:

```
hive> set hive.vectorized.execution.reduce.groupby.enabled;  
hive.vectorized.execution.reduce.groupby.enabled=true  
hive> set hive.vectorized.execution.reduce.enabled;  
hive.vectorized.execution.reduce.enabled=false  
hive> set hive.vectorized.execution.enabled;  
hive.vectorized.execution.enabled=true  
hive>
```

Realizamos 4 ejecuciones activando todos los parámetros de vectorización de Hive:

```
[uamdb20@master ejercicios]$ time hive -f ejercicio6.hql --hiveconf hive.vectorized.execution.enabled=true --hiveconf  
hive.vectorized.execution.reduce.enabled=true --hiveconf hive.vectorized.execution.reduce.groupby.enabled=true
```

```
20674359p.rfihub.com. 4  
2.bp.blogspot.com. 4  
ww.google.es. 2  
as.com. 2  
google.es. 2  
quit.localdomain. 2  
local. 2  
as.es. 2  
quit. 2  
Time taken: 48.167 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release()  
() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real 0m56.751s  
user 0m22.850s  
sys 0m1.760s  
[uamdb20@master ejercicios]$
```

```
2.bp.blogspot.com. 4  
ww.google.es. 2  
as.com. 2  
google.es. 2  
quit.localdomain. 2  
local. 2  
as.es. 2  
quit. 2  
Time taken: 48.49 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFacto  
ry#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an expl  
anation.  
  
real 0m57.161s  
user 0m22.141s  
sys 0m1.800s  
[uamdb20@master ejercicios]$
```

```
1to1.bbva.com. 4  
as.es. 2  
quit. 2  
local. 2  
quit.localdomain. 2  
ww.google.es. 2  
as.com. 2  
google.es. 2  
Time taken: 50.263 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFacto  
ry#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an expl  
anation.  
  
real 0m58.769s  
user 0m21.314s  
sys 0m1.820s  
[uamdb20@master ejercicios]$
```

```
1to1.bbva.com. 4  
as.es. 2  
quit. 2  
local. 2  
quit.localdomain. 2  
ww.google.es. 2  
as.com. 2  
google.es. 2  
Time taken: 51.502 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFacto  
ry#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an expl  
anation.  
  
real 1m0.039s  
user 0m21.357s  
sys 0m1.747s  
[uamdb20@master ejercicios]$
```

Realizamos 4 ejecuciones desactivando todos los parámetros de vectorización de Hive:

```
[uambd20@master ejercicios]$ time hive -f ejercicio6.hql --hiveconf hive.vectorized.execution.enabled=false --hiveconf  
hive.vectorized.execution.reduce.enabled=false --hiveconf hive.vectorized.execution.reduce.groupby.enabled=false
```

```
lto1.bbva.com. 4  
as.es. 2  
quit. 2  
local. 2  
quit.localdomain. 2  
ww.google.es. 2  
as.com. 2  
google.es. 2  
Time taken: 51.606 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real 1m0.026s  
user 0m21.510s  
sys 0m1.679s  
[uambd20@master ejercicios]$
```

```
lto1.bbva.com. 4  
as.es. 2  
quit. 2  
local. 2  
quit.localdomain. 2  
ww.google.es. 2  
as.com. 2  
google.es. 2  
Time taken: 61.243 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real 1m10.104s  
user 0m22.532s  
sys 0m1.637s  
[uambd20@master ejercicios]$
```

```
lto1.bbva.com. 4  
as.es. 2  
quit. 2  
local. 2  
quit.localdomain. 2  
ww.google.es. 2  
as.com. 2  
google.es. 2  
Time taken: 58.123 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real 1m6.712s  
user 0m21.762s  
sys 0m1.824s  
[uambd20@master ejercicios]$
```

```
2.bp.blogspot.com. 4  
ww.google.es. 2  
as.com. 2  
google.es. 2  
quit.localdomain. 2  
local. 2  
as.es. 2  
quit. 2  
Time taken: 50.769 seconds, Fetched: 1024 row(s)  
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.  
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.  
  
real 0m59.210s  
user 0m20.928s  
sys 0m1.452s  
[uambd20@master ejercicios]$
```



Question 10	Mean	Standard Deviation
hive.vectorized.execution.enabled=true hive.vectorized.execution.reduce.enabled=true hive.vectorized.execution.reduce.groupby.enabled=true	58.180	1.312
hive.vectorized.execution.enabled=false hive.vectorized.execution.reduce.enabled=false hive.vectorized.execution.reduce.groupby.enabled=false	64.013	4.565




Figura 3: Resultados de 4 ejecuciones con 3 parámetros diferentes y dos opciones true y false.

RealTime1 indica que `hive.vectorized.execution.enabled=true`,  
`hive.vectorized.execution.reduce.enabled=true`,  
`hive.vectorized.execution.reduce.groupby.enabled=true`.

RealTime2 indica que `hive.vectorized.execution.enabled=false`,  
`hive.vectorized.execution.reduce.enabled=false`,  
`hive.vectorized.execution.reduce.groupby.enabled=false`.

Observamos en la gráfica que al activar todos los parámetros de vectorización de Hive, conseguimos reducir tanto el tiempo de ejecución medio como la desviación estándar del tiempo de ejecución.