

# PRACTICA KAFKA Y LABS

## LAB 1

### Instalar y arrancar Kafka

Descargamos Apache Kafka y lo descomprimos en el directorio de home.

```
[umaster@ibmuamdocke ~]$ tar xvzf ~/Downloads/kafka_2.13-3.7.0.tgz
kafka_2.13-3.7.0/
kafka_2.13-3.7.0/LICENSE
kafka_2.13-3.7.0/NOTICE
kafka_2.13-3.7.0/bin/
kafka_2.13-3.7.0/bin/kafka-delete-records.sh
kafka_2.13-3.7.0/bin/trogdor.sh
```

En el directorio de binarios de Apache Kafka arrancamos Zookeeper.

```
[umaster@ibmuamdocke bin]$ ./zookeeper-server-start.sh ../config/zookeeper.properties
[2024-04-22 13:49:16,849] INFO Reading configuration from: ../config/zookeeper.properties (org.apache.zoo
ookeeper.server.quorum.QuorumPeerConfig)
[2024-04-22 13:49:16,850] WARN ../config/zookeeper.properties is relative. Prepend ./ to indicate that
you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-04-22 13:49:16,853] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.Qu
orumPeerConfig)
[2024-04-22 13:49:16,853] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPe
erConfig)
```

Abrimos otra terminal y, desde el mismo directorio de binarios de Kafka, arrancamos el broker.

```
umaster@ibmuamdocke:~/kafka_2.13-3.7.0/bin x umaster@ibmuamdocke:~/kafka_2.13-3.7.0/bin x
[umaster@ibmuamdocke bin]$ ./kafka-server-start.sh ../config/server.properties
[2024-04-22 13:51:37,641] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log4jCont
rollerRegistration)
[2024-04-22 13:51:37,856] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disable cl
ient-initiated TLS renegotiation (org.apache.zookeeper.common.X509Util)
[2024-04-22 13:51:37,923] INFO Registered signal handlers for TERM, INT, HUP (org.apache.kafka.common.u
tils.LoggingSignalHandler)
[2024-04-22 13:51:37,924] INFO starting (kafka.server.KafkaServer)
[2024-04-22 13:51:37,924] INFO Connecting to zookeeper on localhost:2181 (kafka.server.KafkaServer)
[2024-04-22 13:51:37,933] INFO [ZooKeeperClient Kafka server] Initializing a new session to localhost:2
181. (kafka.zookeeper.ZooKeeperClient)
```

## Productor y Consumidor

Creamos un Topic con Productor y Consumidor y lo validamos.  
(Ejecuto el comando con la ruta completa porque no funciona con bin/kafka-topics.sh.)

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-topics.sh --bootstrap-server local
host:9092 --replication-factor 1 --partitions 1 --create --topic test
Created topic test.
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-topics.sh --bootstrap-server local
host:9092 --list
test
[umaster@ibmuamdocke bin]$
```

Producimos y consumimos mensajes

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-console-producer.sh --brok
er-list localhost:9092 --topic test
>message00
>message01
>^[[A^[[Bmessage02
```

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-console-consumer.sh --boot
strap-server localhost:9092 --topic test --from-beginning
message00
message01
message02
^CProcessed a total of 3 messages
```

## Almacenamiento de los mensajes

Revisamos donde están almacenados los mensajes:

```
[umaster@ibmuamdocke bin]$ tree /tmp/kafka-logs
/tmp/kafka-logs
├── cleaner-offset-checkpoint
├── consumer_offsets-0
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── consumer_offsets-1
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── consumer_offsets-10
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── consumer_offsets-11
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── consumer_offsets-12
│   ├── 00000000000000000000.index
│   └── 00000000000000000000.log
```

```

├── consumer_offsets-8
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── consumer_offsets-9
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata
├── log-start-offset-checkpoint
├── meta.properties
├── recovery-point-offset-checkpoint
├── replication-offset-checkpoint
├── test-0
│   ├── 00000000000000000000.index
│   ├── 00000000000000000000.log
│   ├── 00000000000000000000.timeindex
│   ├── leader-epoch-checkpoint
│   └── partition.metadata

```

51 directories, 260 files

[umaster@ibmuamdocker bin]\$ █

```

[umaster@ibmuamdocker bin]$ cd /tmp/kafka-logs/test-0/
[umaster@ibmuamdocker test-0]$ od -bc 00000000000000000000.log | more
00000000 000 000 000 000 000 000 000 000 000 000 000 101 000 000 000 000
          \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 A \0 \0 \0 \0
00000020 002 044 125 245 112 000 000 000 000 000 000 000 000 001 217 005
          002 $ U 245 J \0 \0 \0 \0 \0 \0 \0 \0 \0 001 217 005
00000040 343 344 244 000 000 001 217 005 343 344 244 000 000 000 000 000
          343 344 244 \0 \0 001 217 005 343 344 244 \0 \0 \0 \0 \0
00000060 000 000 000 000 000 000 000 000 000 000 000 000 001 036 000 000
          \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 001 036 \0 \0
00000100 000 001 022 155 145 163 163 141 147 145 060 060 000 000 000 000
          \0 001 022 m e s s a g e 0 0 \0 \0 \0 \0
00000120 000 000 000 000 001 000 000 000 101 000 000 000 000 002 070 176
          \0 \0 \0 \0 001 \0 \0 \0 A \0 \0 \0 \0 002 8 ~
00000140 160 156 000 000 000 000 000 000 000 000 001 217 005 343 375 052
          p n \0 \0 \0 \0 \0 \0 \0 \0 \0 001 217 005 343 375 *
00000160 000 000 001 217 005 343 375 052 000 000 000 000 000 000 000 000
          \0 \0 001 217 005 343 375 * \0 \0 \0 \0 \0 \0 \0 \0
00000200 000 000 000 000 000 001 000 000 000 001 036 000 000 000 001 022

```

```

0000200 000 000 000 000 000 001 000 000 000 001 036 000 000 000 001 022
          \0 \0 \0 \0 \0 001 \0 \0 \0 001 036 \0 \0 \0 \0 001 022
0000220 155 145 163 163 141 147 145 060 061 000 000 000 000 000 000 000
          m e s s a g e 0 1 \0 \0 \0 \0 \0 \0 \0
0000240 000 002 000 000 000 107 000 000 000 000 002 317 374 357 150 000
          \0 002 \0 \0 \0 G \0 \0 \0 \0 002 317 374 357 h \0
0000260 000 000 000 000 000 000 000 001 217 005 344 035 213 000 000 001
          \0 \0 \0 \0 \0 \0 \0 001 217 005 344 035 213 \0 \0 001
0000300 217 005 344 035 213 000 000 000 000 000 000 000 000 000 000 000
          217 005 344 035 213 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0
0000320 000 000 002 000 000 000 001 052 000 000 000 001 036 033 133 101
          \0 \0 002 \0 \0 \0 001 * \0 \0 \0 001 036 033 [ A
0000340 033 133 102 155 145 163 163 141 147 145 060 062 000

```

--More--

## Borrar Topics

Borramos el topic y arrancamos el bróker y comprobamos los topics disponibles para validar.

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-topics.sh --bootstrap-server localhost:9092 --delete --topic test
```

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
__consumer_offsets
```

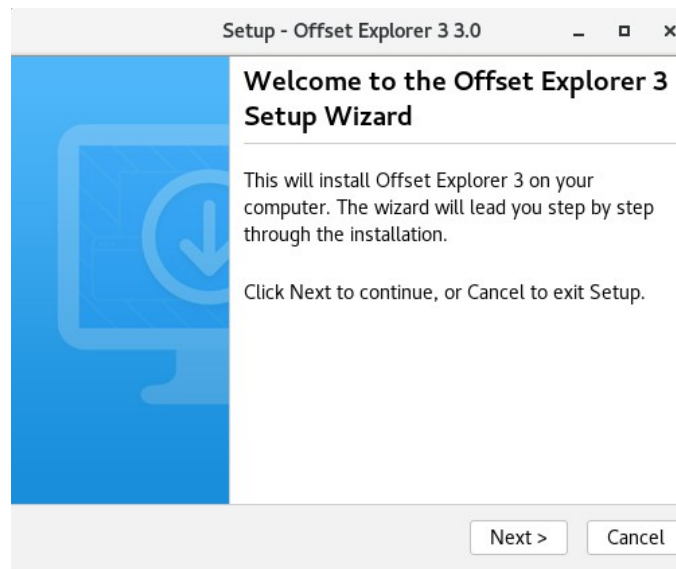
Observamos que no existe un topic llamado "test" disponible.

## LAB 2

### Instalar Offset Explorer

Descargamos Offset Explorer y lo ejecutamos:

```
[umaster@ibmuamdocke ~]$ sh ~/Downloads/offsetexplorer.sh  
Starting Installer ...
```

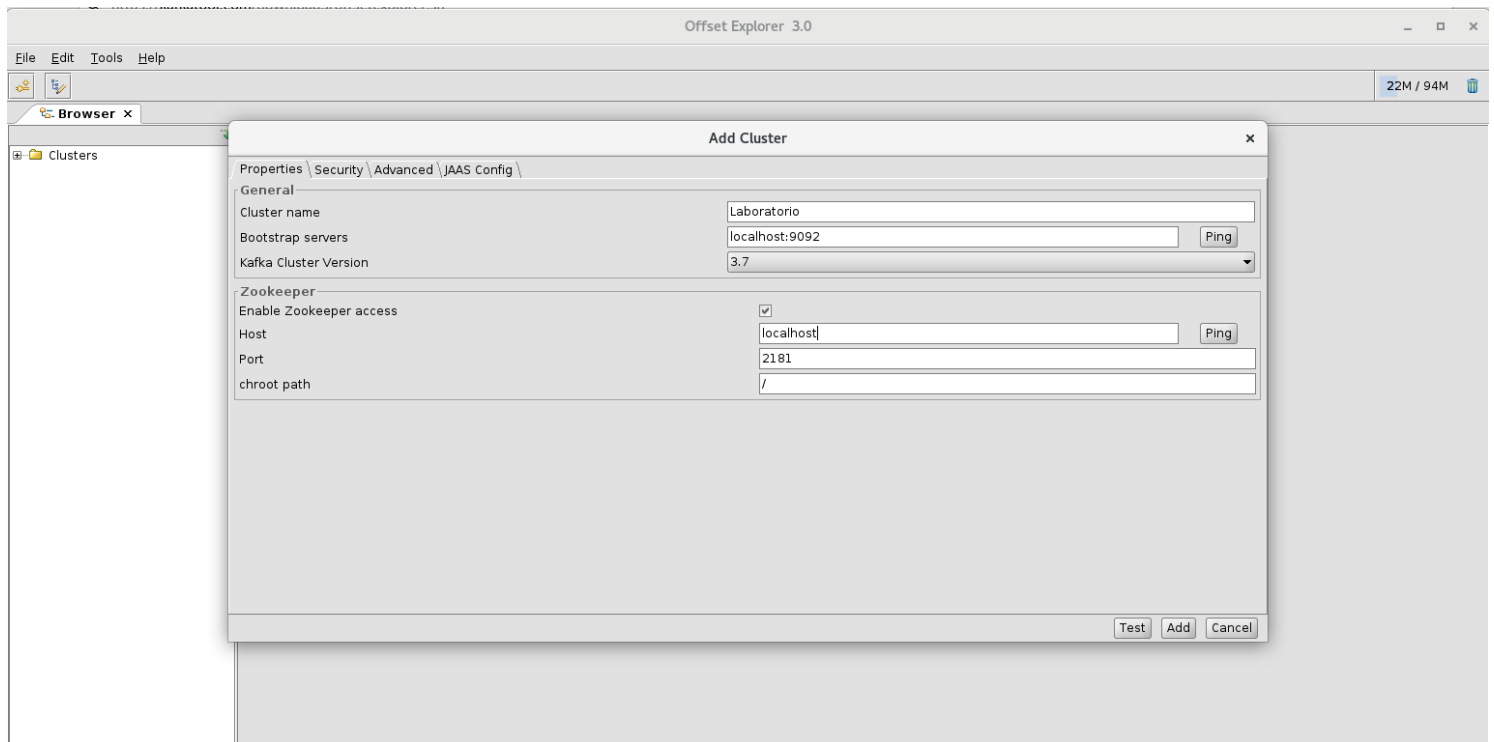


Hemos realizado la instalación, aceptando acuerdo de licencia, directorio por defecto de la instalación, directorio para la creación de los enlaces simbólicos.

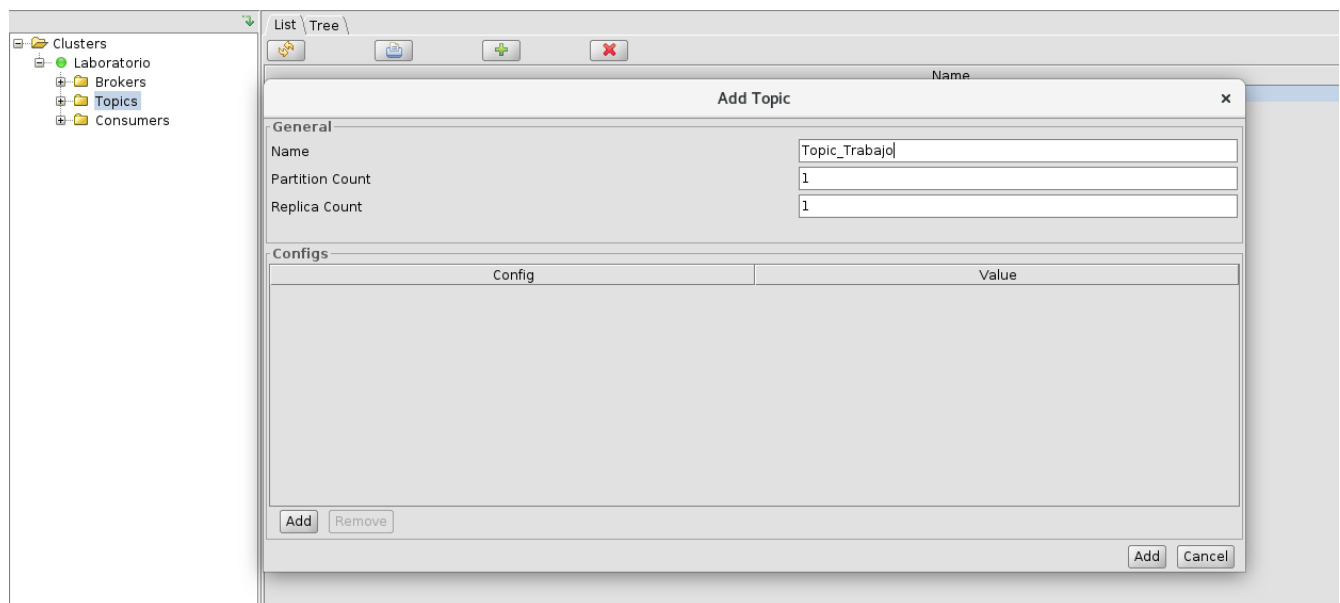
### Configurar y comprobar

```
[umaster@ibmuamdocke ~]$ cd ~/offsetexplorer3  
[umaster@ibmuamdocke offsetexplorer3]$ ./offsetexplorer
```

## Creamos una conexión a nuestro cluster



## Añadimos un topic



Hemos establecido que la clave y el valor se representarán como un “Byte Array”.

Properties \ Data \ Partitions \ Config \

General

Topic Name

Topic\_Trabajo

Content Types

Key

Byte Array

Value

Byte Array

\* Changing the type affects how messages are shown/added in the partition Data panel

Messages

Total number of messages

Click 'Refresh' to get count

Refresh

Añadimos un mensaje.  
Observamos los mensajes del topic:

Clusters

Laboratorio

Brokers

Topics

\_consumer\_

Topic\_Trabaj

Partitions

Partiti

Consumers

console-con

Properties \ Data \ Partitions \ Config \

Filter

Messages Oldest

Partition	Offset	Key	Value	Timestamp
0	0	4B657931	457320756E20707275...	2024-04-22 18:26:02.104

Partition \ Offset \ Key \ Value \ Timestamp \ Headers \

Headers

View As Text

Key	Value
Key1	Es un prueba

Ready [Messages = 1] [12 Bytes] [105 ms]

Max Messages (per partition) 50

## LAB 3

### Usando Kafka Connect

Editamos el fichero config/connect-standalone.properties

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ cd config/  
[umaster@ibmuamdock config]$ nano connect-standalone.properties  
[umaster@ibmuamdock config]$
```

GNU nano 2.3.1	File: connect-standalone.properties	Modified
<pre># (connectors, converters, transformations). The list should consist of to\$ # any combination of: # a) directories immediately containing jars with plugins and their depend\$ # b) uber-jars with plugins and their dependencies # c) directories immediately containing the package directory structure of\$ # Note: symlinks will be followed to discover dependencies or plugins. # Examples: # plugin.path=/usr/local/share/java,/usr/local/share/kafka/plugins,/opt/co\$ #plugin.path= plugin.path=/home/umaster/kafka_2.13-3.7.0/libs/connect-file-3.7.0.jar</pre>		

Creamos el fichero test.txt e insertamos algunos valores

```
[umaster@ibmuamdock config]$ cd ~/kafka_2.13-3.7.0/bin  
[umaster@ibmuamdock bin]$ touch test.txt  
[umaster@ibmuamdock bin]$ echo "Hola" >> test.txt  
[umaster@ibmuamdock bin]$ echo "Que pasa" >> test.txt
```

Observamos que el archivo test.txt existe y contiene las líneas: “Hola” y “Que pasa”

```
[umaster@ibmuamdock bin]$ cat test.txt  
Hola  
Que pasa
```

Arrancamos los conectores

```
[umaster@ibmuamdock bin]$ /home/umaster/kafka_2.13-3.7.0/bin/connect-standalone.sh /home/umaster/kafka_2.13-3.7.0/config/  
connect-standalone.properties /home/umaster/kafka_2.13-3.7.0/config/connect-file-source.properties /home/umaster/kafka_2.13  
-3.7.0/config/connect-file-sink.properties
```



Comprobamos el contenido del fichero test.sink.txt en el directorio bin

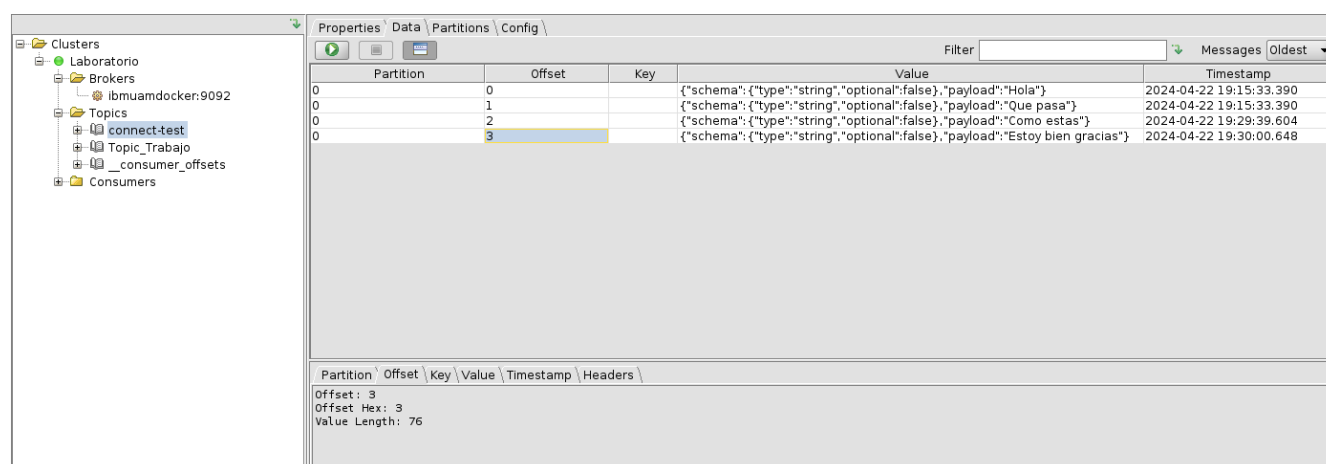
```
[umaster@ibmuamdocke bin]$ cat test.sink.txt
Hola
Que pasa
```

Comprobamos el contenido del topic que se ha creado

```
[umaster@ibmuamdocke bin]$ /home/umaster/kafka_2.13-3.7.0/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic connect-test --from-beginning
{"schema":{"type":"string","optional":false},"payload":"Hola"}
{"schema":{"type":"string","optional":false},"payload":"Que pasa"}
```

Incluimos nuevas entradas en el fichero test.txt

```
[umaster@ibmuamdocke bin]$ touch test.txt
[umaster@ibmuamdocke bin]$ echo "Como estas" >> test.txt
[umaster@ibmuamdocke bin]$ echo "Estoy bien gracias" >> test.txt
[umaster@ibmuamdocke bin]$ cat test.txt
Hola
Que pasa
Como estas
Estoy bien gracias
[umaster@ibmuamdocke bin]$ cat test.sink.txt
Hola
Que pasa
Como estas
Estoy bien gracias
```



Partition	Offset	Key	Value	Timestamp
0	0		{"schema":{"type":"string","optional":false},"payload":"Hola"}	2024-04-22 19:15:33.390
0	1		{"schema":{"type":"string","optional":false},"payload":"Que pasa"}	2024-04-22 19:15:33.390
0	2		{"schema":{"type":"string","optional":false},"payload":"Como estas"}	2024-04-22 19:29:39.604
0	3		{"schema":{"type":"string","optional":false},"payload":"Estoy bien gracias"}	2024-04-22 19:30:00.648

Partition: 0  
Offset: 3  
Offset Hex: 3  
Value Length: 76

Hemos verificado que los mensajes se incluyen correctamente al topic “connect-test” y al fichero de destino test.sink.txt.

## LAB 4

### Desarrollo de clientes con python

Hemos actualizado pip e instalado los módulos necesarios

```
[umaster@ibmuamdocke ~]$ pip3 install --upgrade pip --user
Requirement already satisfied: pip in /usr/local/lib/python3.6/site-packages (21.3.1)
[umaster@ibmuamdocke ~]$ pip3 install kafka-Python --user
Requirement already satisfied: kafka-Python in ./local/lib/python3.6/site-packages (2.0.2)
```

Adjunto los archivos de log llamados consumer.log y producer.log a esta práctica.

### Plantilla de Productor

Hemos rellenado los codigos.

```
from kafka import KafkaProducer
from faker import Faker
from json import dumps, loads
import time
import logging
import sys
import random
```

```
: #Logging
logging.basicConfig(format='%(asctime)s %(message)s',
                    datefmt='%d-%m-%Y %H:%M:%S',
                    filename='producer.log',
                    filemode='w')

logger = logging.getLogger()
logger.setLevel(logging.INFO)
```

```
: #Faker
fake=Faker('es_ES')
```

```
#Configuración del entorno Kafka
topic_name = 'my-topic'
key = 'my-key'
bootstrap_servers = ['localhost:9092', 'localhost:9093']
logging.info('Arrancando...')
```

```
: #Definiendo el conector
def connect_kafka_producer(bootstrap_servers):
    producer = None
    try:
        producer = KafkaProducer(bootstrap_servers=bootstrap_servers,
                                key_serializer=str.encode(),
                                value_serializer=lambda x: dumps(x).encode('utf-8'))
    except Exception as ex:
        logging.error('Exception while connecting producer')
        logging.error(ex)
    finally:
        return producer
```

```

: #Definiendo la publicación del mensaje
def publish_message(producer, topic, key, value):
    try:
        producer.send(topic, key=key, value=value)
        producer.flush()
        logging.debug('Publicación del Mensaje OK')
    except Exception as ex:
        logging.error(value)
        logging.error('Excepción al publicar el mensaje')
        logging.error(ex)
    finally:
        return

```

```

producer = connect_kafka_producer(bootstrap_servers)
logging.info("Conexión a broker: "+str(producer))
for i in range(30):
    data={
        'user_id': fake.random_int(min=20000, max=100000),
        'user_name':fake.name(),
        'user_address':fake.street_address() + ' | ' + fake.city() + ' | ' + fake.country_code(),
        'platform': random.choice(['Mobile', 'Laptop', 'Tablet']),
        'signup_at': str(fake.date_time_this_month())
    }
    publish_message(producer, topic_name, key, data)
    logging.info(data)
    time.sleep(1)

```

## Plantilla de Consumidor

```
from kafka import KafkaConsumer
from json import loads
import logging
```

```
#Logging
logging.basicConfig(format='%(asctime)s %(message)s',
                    datefmt='%d-%m-%Y %H:%M:%S',
                    filename='consumer.log',
                    filemode='w')

logger = logging.getLogger()
logger.setLevel(logging.INFO)
```

```
#Configuración del entorno Kafka
topic_name = 'my-topic'
key = 'my-key'
bootstrap_servers = ['localhost:9092']
logging.info('Arrancando...')
```

```
try:
    consumer = None
    consumer = KafkaConsumer(topic_name,
                             bootstrap_servers=bootstrap_servers,
                             auto_offset_reset='earliest',
                             enable_auto_commit=True,
                             group_id='my-processing-group',
                             value_deserializer=lambda x: loads(x.decode('utf-8')))
except Exception as ex:
    logging.error('Exception while connecting consumer')
    logging.error(ex)
```

```
for message in consumer:
    message = message.value
    print('Leido: {}'.format(message))
    logger = logging.info('Mensaje: {}'.format(message))
```

## LAB 5

### Configuración y arranque de nuevos brokers

Copiamos las propiedades en el directorio de config en el directorio Apache Kafka

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ cp config/server.properties config/server1.properties
[umaster@ibmuamdock kafka_2.13-3.7.0]$ cp config/server.properties config/server2.properties
```

En server1.properties establecemos los siguientes valores:

```
# The id of the broker. This must be set to a unique integer for each broker.
broker.id=1

##### Socket Server Settings #####

# The address the socket server listens on. If not configured, the host name will be equal to the value of
# java.net.InetAddress.getCanonicalHostName(), with PLAINTEXT listener name, and port 9092.
#   FORMAT:
#     listeners = listener_name://host_name:port
#   EXAMPLE:
#     listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9093

# A comma separated list of directories under which to store log files
log.dirs=/tmp/kafka-logs-1
```

En server2.properties establecemos los siguientes valores:

```
# The id of the broker. This must be set to a unique integer for each broker.
broker.id=2

##### Socket Server Settings #####

# The address the socket server listens on. If not configured, the host name will be equal to the value of
# java.net.InetAddress.getCanonicalHostName(), with PLAINTEXT listener name, and port 9092.
#   FORMAT:
#     listeners = listener_name://host_name:port
#   EXAMPLE:
#     listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9094

# A comma separated list of directories under which to store log files
log.dirs=/tmp/kafka-logs-2
```

## Arrancamos todos los brokers

```
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ nohup bin/kafka-server-start.sh config/server.properties > /dev/null 2>&1 &
[1] 14148
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ nohup bin/kafka-server-start.sh config/server1.properties > /dev/null 2>&1 &
[2] 14584
[1] Exit 1 nohup bin/kafka-server-start.sh config/server.properties > /dev/null 2>&1
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ nohup bin/kafka-server-start.sh config/server2.properties > /dev/null 2>&1 &
[3] 15010
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ █
```

## Creación de diferentes topics

Creamos topics con diferentes valores de factor de replicación y de particiones

Factor de replica 3 | Particiones 1 (nombre del topic my-rep3-part1-topic)

```
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --replication-factor 3 --partitions 1 --create --topic my-rep3-part1-topic
Created topic my-rep3-part1-topic.
```

```
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic my-rep3-part1-topic
Topic: my-rep3-part1-topic      TopicId: 6hZkG62jRrq083jKHGrGrQ PartitionCount: 1      ReplicationFactor: 3   Configs:
      Topic: my-rep3-part1-topic Partition: 0    Leader: 1      Replicas: 1,0,2 Isr: 1,0,2
[umaster@ibmuamdocker kafka_2.13-3.7.0]$
```

Podemos observar nombre de topic, topic id, factor de replicación y particiones.

Repetimos para:

Factor de replica 3 | Particiones 3 (nombre del topic my-rep3-part3-topic)

```
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --replication-factor 3 --partitions 3 --create --topic my-rep3-part3-topic
Created topic my-rep3-part3-topic.
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic my-rep3-part3-topic
Topic: my-rep3-part3-topic      TopicId: 7GeTYE7fTJ0X75_0290v4A PartitionCount: 3      ReplicationFactor: 3   Configs:
      Topic: my-rep3-part3-topic Partition: 0    Leader: 1      Replicas: 1,2,0 Isr: 1,2,0
      Topic: my-rep3-part3-topic Partition: 1    Leader: 0      Replicas: 0,1,2 Isr: 0,1,2
      Topic: my-rep3-part3-topic Partition: 2    Leader: 2      Replicas: 2,0,1 Isr: 2,0,1
```

Tenemos 3 particiones y 3 replicas.

```
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --replication-factor 1 --partitions 3 --create --topic my-repl-part3-topic
Created topic my-repl-part3-topic.
[umaster@ibmuamdocker kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic my-repl-part3-topic
Topic: my-repl-part3-topic      TopicId: kf-zPl4JQdKrKqX5dUMfg PartitionCount: 3      ReplicationFactor: 1 Configs:
  Topic: my-repl-part3-topic    Partition: 0    Leader: 1        Replicas: 1      Isr: 1
  Topic: my-repl-part3-topic    Partition: 1    Leader: 0        Replicas: 0      Isr: 0
  Topic: my-repl-part3-topic    Partition: 2    Leader: 2        Replicas: 2      Isr: 2
```

Validamos “logs”

```
$ tree /tmp/kafka-logs
```

```
/tmp/kafka-logs  
├── cleaner-offset-checkpoint  
├── connect-test-0  
│   ├── 00000000000000000000000000000000.index  
│   ├── 00000000000000000000000000000000.log  
│   ├── 00000000000000000000000000000000.timeindex  
│   ├── leader-epoch-checkpoint  
│   └── partition.metadata  
├── consumer_offsets-0  
│   ├── 00000000000000000000000000000000.index  
│   ├── 00000000000000000000000000000000.log  
│   ├── 00000000000000000000000000000000.timeindex  
│   ├── leader-epoch-checkpoint  
│   └── partition.metadata  
├── consumer_offsets-1  
│   ├── 00000000000000000000000000000000.index  
│   ├── 00000000000000000000000000000000.log  
│   ├── 00000000000000000000000000000000.timeindex  
│   ├── leader-epoch-checkpoint  
│   └── partition.metadata
```

```
└─ recovery-point-offset-checkpoint
└─ replication-offset-checkpoint
└─ Topic_Trabajo-0
    └─ 00000000000000000000000000000000.index
    └─ 00000000000000000000000000000000.log
    └─ 00000000000000000000000000000000.timeindex
    └─ leader-epoch-checkpoint
    └─ partition.metadata
```

```
57 directories, 290 files
```

**\$ tree /tmp/kafka-logs-1**

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ tree /tmp/kafka-logs-1
/tmp/kafka-logs-1
├── cleaner-offset-checkpoint
├── log-start-offset-checkpoint
├── meta.properties
├── recovery-point-offset-checkpoint
└── replication-offset-checkpoint

0 directories, 5 files
```

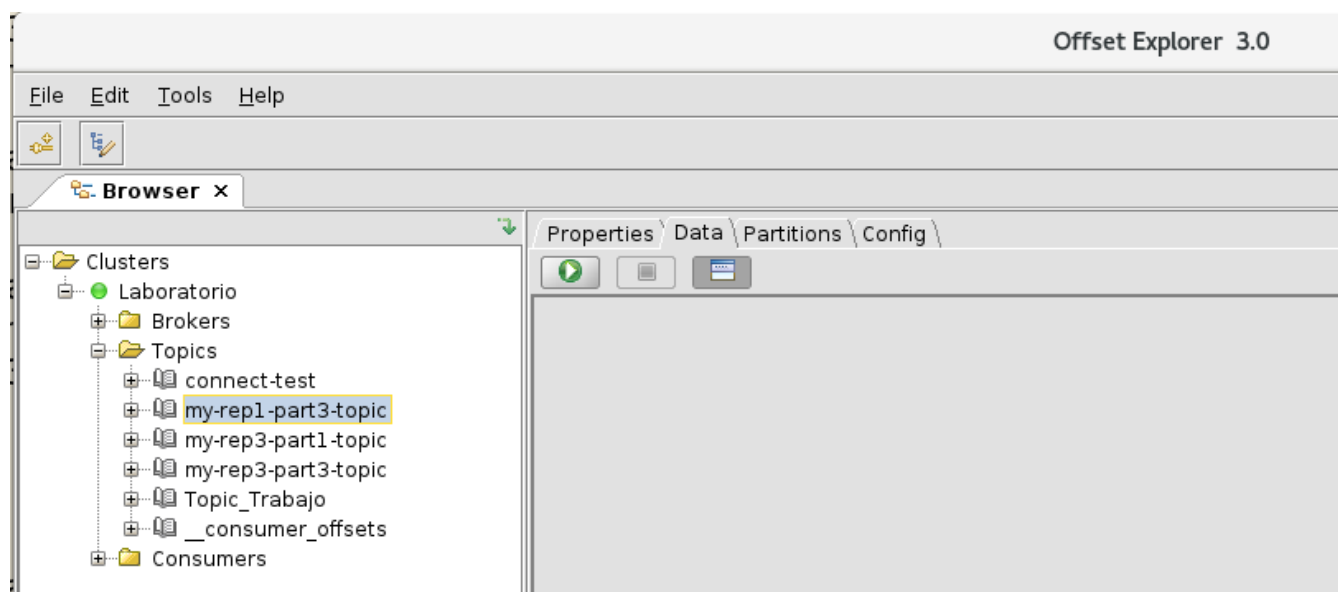
**\$ tree /tmp/kafka-logs-2**

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ tree /tmp/kafka-logs-2
/tmp/kafka-logs-2
├── cleaner-offset-checkpoint
├── log-start-offset-checkpoint
├── meta.properties
├── recovery-point-offset-checkpoint
└── replication-offset-checkpoint

0 directories, 5 files
[umaster@ibmuamdock kafka_2.13-3.7.0]$
```

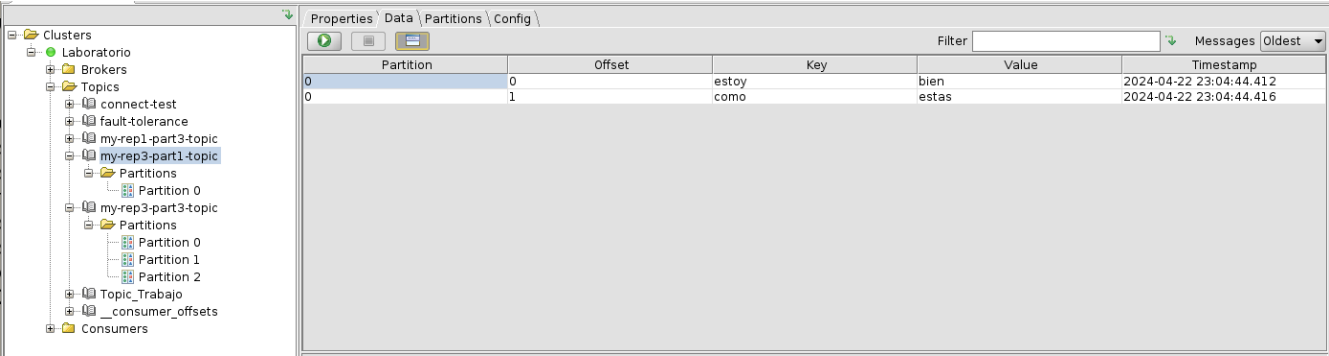
Observamos directorios de registro en el directorio kafka-logs, pero no existen directorios de registro en los directorios kafka-logs-1 y kafka-logs-2.

Vemos que creamos los topics y podemos observar en Offset Explorer:



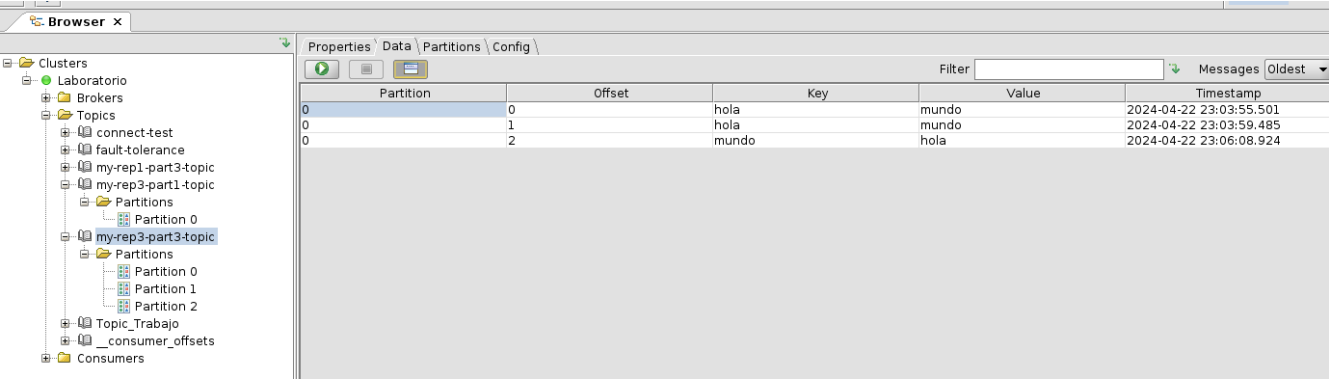


Introducimos valores en los topics y comprobamos resultados con el Offset Explorer:



The screenshot shows the Kafka Offset Explorer interface. On the left, a tree view displays the cluster structure, with 'my-rep3-part1-topic' selected. The main pane shows the 'Data' tab for this topic, displaying a table of messages. The table has columns for Partition, Offset, Key, Value, and Timestamp. Two messages are visible: one with offset 0 and key 'estoy' (value 'bien'), and another with offset 1 and key 'como' (value 'estas').

Partition	Offset	Key	Value	Timestamp
0	0	estoy	bien	2024-04-22 23:04:44.412
0	1	como	estas	2024-04-22 23:04:44.416



The screenshot shows the Kafka Offset Explorer interface. On the left, a tree view displays the cluster structure, with 'my-rep3-part3-topic' selected. The main pane shows the 'Data' tab for this topic, displaying a table of messages. The table has columns for Partition, Offset, Key, Value, and Timestamp. Three messages are visible: one with offset 0 and key 'hola' (value 'mundo'), one with offset 1 and key 'hola' (value 'mundo'), and one with offset 2 and key 'mundo' (value 'hola').

Partition	Offset	Key	Value	Timestamp
0	0	hola	mundo	2024-04-22 23:03:55.501
0	1	hola	mundo	2024-04-22 23:03:59.485
0	2	mundo	hola	2024-04-22 23:06:08.924

Al añadir nuevos valores, se pueden observar los pares clave-valor en los topics.

## Multibroker. Comprobar tolerancia a fallos

Obtenemos el PID con el comando “ps”

**\$ ps -ef | grep server**

PID: 2951, 3379, 3405, 4408, 4479, 9017, 9665, 14584, 15010

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ ps -ef | grep server
umaster 2951 1 0 12:38 ? 00:00:00 /usr/libexec/gnome-shell-calendar-server
umaster 3379 3235 0 12:38 ? 00:00:00 /usr/libexec/evolution-calendar-factory-subprocess --factory all --bus-name org.gnome.evolution.data-server.Subprocess.Backend.Calendarx3235x2 --own-path /org/gnome/evolution/data-server/Subprocess/Backend/Calendar/3235/2
umaster 3405 3389 0 12:38 ? 00:00:00 /usr/libexec/evolution-addressbook-factory-subprocess --factory all --bus-name org.gnome.evolution.data-server.Subprocess.Backend.AddressBookx3389x2 --own-path /org/gnome/evolution/data-server/Subprocess/Backend/AddressBook/3389/2
umaster 4408 1 0 12:42 ? 00:00:18 /usr/libexec/gnome-terminal-server
umaster 4479 30779 0 19:15 pts/3 00:00:26 java -Xms256M -Xmx2G -server -XX:+UseG1GC -XX:Max
```

```
umaster 9017 4415 0 13:49 pts/0 00:00:18 java -Xmx512M -Xms512M -server
```

```
umaster 9665 9575 0 13:51 pts/1 00:02:11 java -Xmx1G -Xms1G -server
```

```
umaster 14584 1 0 21:12 ? 00:00:12 java -Xmx1G -Xms1G -server
```

```
umaster 15010 1 0 21:12 ? 00:00:12 java -Xmx1G -Xms1G -server
```

**\$ ps -ef | grep server1**

PID: 14584

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ ps -ef | grep server1
umaster 14584 1 0 21:12 ? 00:00:12 java -Xmx1G -Xms1G -server -XX:+UseG1GC -XX:Max
GCPauseMillis=20 -XX:InitiatingHeapOccupancyPercent=35 -XX:+ExplicitGCInvokesConcurrent -XX:Max
```

**\$ ps -ef | grep server2**

PID: 15010

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ ps -ef | grep server2
umaster 15010      1  0 21:12 ?                00:00:12 java -Xmx1G -Xms1G -server -XX:+UseG1GC
```

Hemos matado el proceso del broker 2 y hemos verificado mediante el comando ps que se ha matado correctamente.

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ kill 15010
```

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ ps -ef | grep server2
umaster 24471 22110  0 22:40 pts/6      00:00:00 grep --color=auto server2
```

Creamos un nuevo topic y revisamos la configuración.

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092
--replication-factor 2 --partitions 2 --create --topic fault-tolerance
Created topic fault-tolerance.
[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092
--describe --topic fault-tolerance
Topic: fault-tolerance TopicId: Gi5q84ttRJ2U--h4ciDCdA PartitionCount: 2      ReplicationFact
or: 2  Configs:
      Topic: fault-tolerance Partition: 0    Leader: 1      Replicas: 1,0  Isr: 1,0
      Topic: fault-tolerance Partition: 1    Leader: 0      Replicas: 0,1  Isr: 0,1
[umaster@ibmuamdock kafka_2.13-3.7.0]$
```

Observamos que hemos creado un topic y tiene 2 particiones y 2 réplicas con el nombre del topic.

Producimos y consumimos nuevos mensajes

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-console-producer.sh --broker-list localhost:
9092, localhost:9093 --topic fault-tolerance
>XX
>YY
-
^C[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-console-consumer.sh --bootstrap-server lo
host:9092, localhost:9093 --from-beginning --topic fault-tolerance
XX
YY
^CProcessed a total of 2 messages
[umaster@ibmuamdock kafka_2.13-3.7.0]$
```

Sabemos que PID del broker 1 es 14584

Hemos matado el proceso del broker 1 y hemos verificado mediante el comando ps que se ha matado correctamente.

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ kill 14584
[umaster@ibmuamdock kafka_2.13-3.7.0]$ ps -ef | grep server1
umaster 27015 22110 0 22:54 pts/6 00:00:00 grep --color=auto server1
[umaster@ibmuamdock kafka_2.13-3.7.0]$
```

Revisamos la configuración del topic y validamos que seguimos disponiendo de los mensajes:

```
[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-topics.sh --bootstrap-server localhost:9092
--describe --topic fault-tolerance
Topic: fault-tolerance TopicId: Gi5q84ttRJ2U--h4ciDCdA PartitionCount: 2 ReplicationFactor: 2
Configs:
  Topic: fault-tolerance Partition: 0 Leader: 0 Replicas: 1,0 Isr: 0
  Topic: fault-tolerance Partition: 1 Leader: 0 Replicas: 0,1 Isr: 0
[umaster@ibmuamdock kafka_2.13-3.7.0]$ bin/kafka-console-consumer.sh --bootstrap-server localhost:9092, localhost:9093 --from-beginning --topic fault-tolerance
XX
YY
^CProcessed a total of 2 messages
[umaster@ibmuamdock kafka_2.13-3.7.0]$
```

En Offset Explorer vemos que tenemos todos los topics que hemos creado